


One official test run: 960 samples, fault injected at sample 161



The recall of every fault class is capped at $800/960 = 0.8333$ under literal labels. Across seven classifiers the relative accuracy gain is 0.142-0.150, against a structural bound of 0.1587: the effect is the convention, not the model.